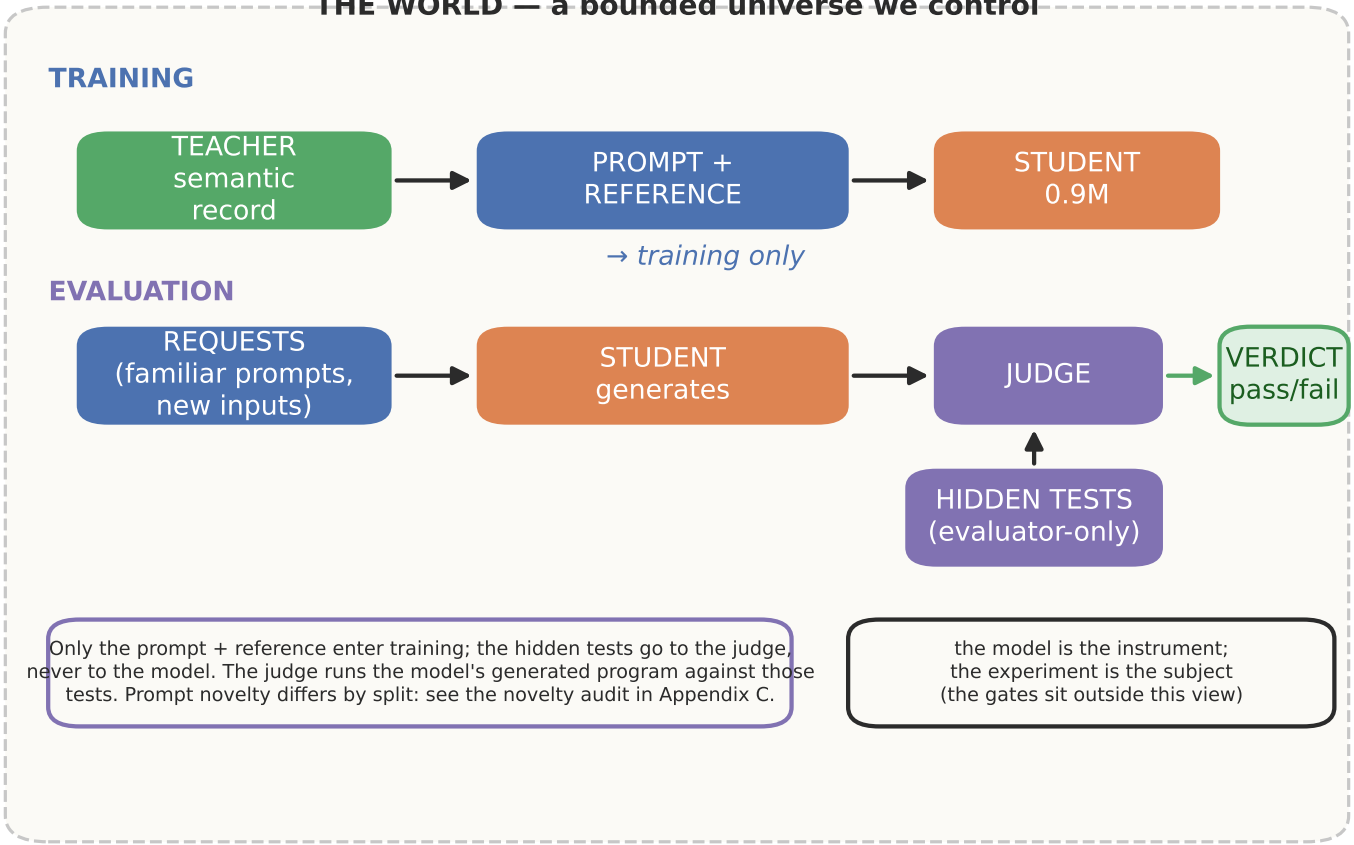
THE WORLD — a bounded universe we control
TRAINING
TEACHER
semantic
record
PROMPT +
REFERENCE
STUDENT
0.9M
→ training only
EVALUATION
REQUESTS
(familiar prompts,
new inputs)
STUDENT
generates
JUDGE
VERDICT
pass/fail
HIDDEN TESTS
(evaluator-only)
Only the prompt + reference enter training; the hidden tests go to the judge, never to the model. The judge runs the model's generated program against those tests. Prompt novelty differs by split: see the novelty audit in Appendix C.
the model is the instrument;
the experiment is the subject
(the gates sit outside this view)