THE BREW LOOP

dev feedback: what we measure (loss / dev score) shapes the next round of data & hyperparams.

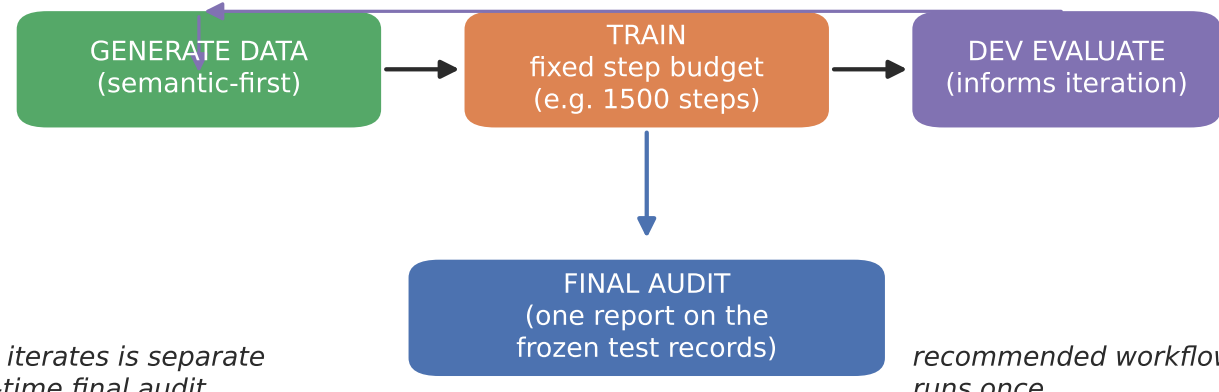

the dev loop that iterates is separate from the one-time final audit.

each dev pass feeds the next: what we measure shapes what we generate; but the headline numbers come from a single frozen audit, not an adaptive one.

recommended workflow: the final audit runs once on the frozen records and is NOT fed back to change training. (In our runs the pilot test records reused the training prompt strings -- see the novelty audit -- so 'frozen' refers to the record list, not to genuinely-seen-new prompts.)