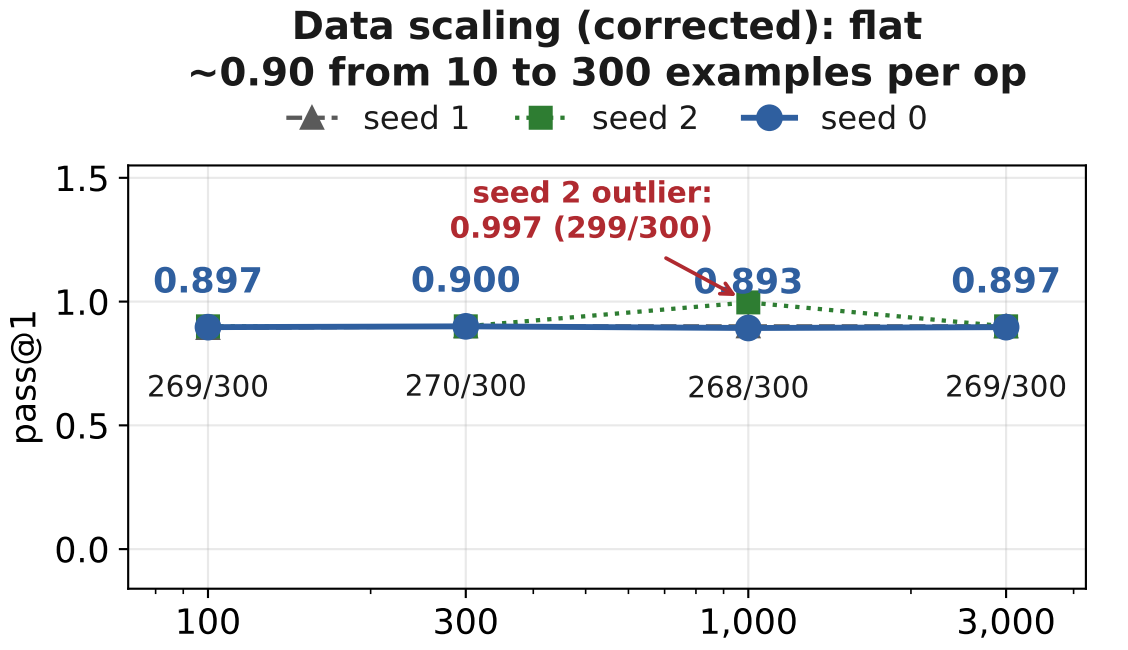

labels/counts = seed 0 (successes / 300 shared test, 30/op x 10 ops)
one fixed tokenizer · 1,000 steps · 3 seeds · all 12 arms cover 10/10 ops
template 0.800; retrieval 0.943-1.000 in all 12 arms
per seed (100/300/1k/3k): s0 .897/.900/.893/.897 s1 .893/.900/.900/.900
s2 .900/.900/.997/.900 -- the 0.997 (size 1,000) is the marked outlier
cond_abs_max is 0.0 in 9/12 arms; 0.100->0.641 = flawed run (P0-01)