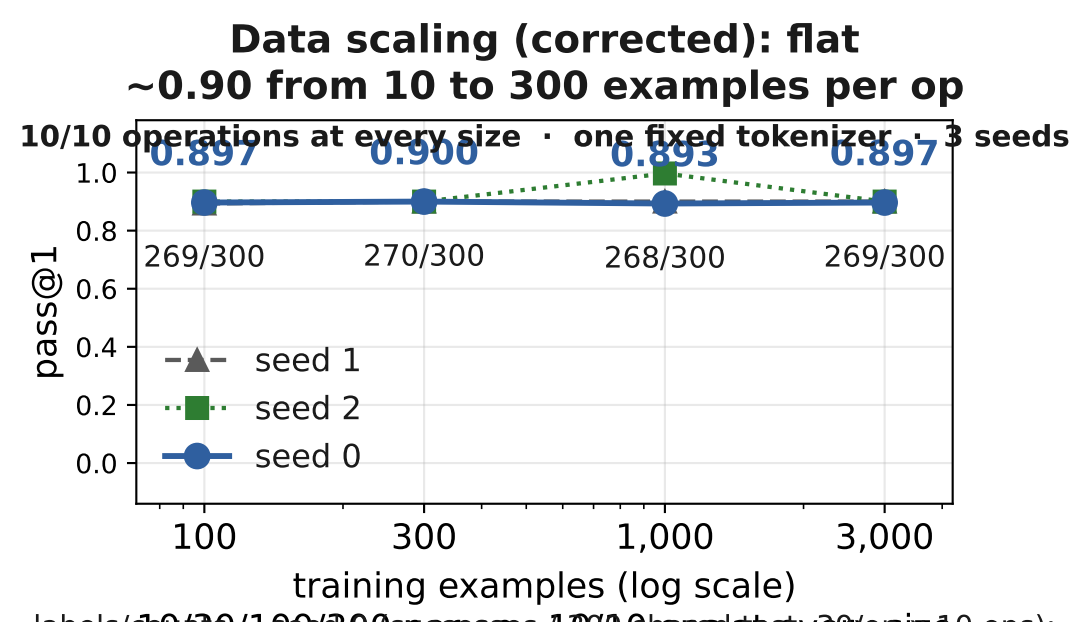

labels/chunk=1 best loss across 1,000 steps; test = 20 (size 10 ops);
template 0.800 and retrieval 0.943-1.000 in all 12 arms; 1,000 steps;
the earlier 0.100->0.641 curve is a flawed historical run (P0-01)