THE JUDGE: design stages vs what is recorded

LAYER 1 — DESIGN: four diagnostic stages (desired)

which stage failed was NOT saved per candidate

1 parse → **2 AST allowlist** → **3 execute** → **4 semantic**

ideal taxonomy — NOT recorded per candidate

LAYER 2 — RECORDED: combined gate + 2 counters

COMBINED VALIDATION GATE — sandbox.validate_source() parses AND checks the allowlist in ONE call → sets syntax_ok & capability_ok together

→ validation_pass (any fail → score 0)
'syntax' rate includes capability rejects

EXECUTE program against the hidden tests (loop)

tests_passed ≥ 1
→ 'executable'
(at_least_one_test_pass)

ALL tests pass
→ pass@1 = 1
(all_tests_pass)

'*executable*' = *passed ≥ 1 hidden test, not merely 'ran'*

per-candidate failure taxonomy
(which stage failed) was NOT collected