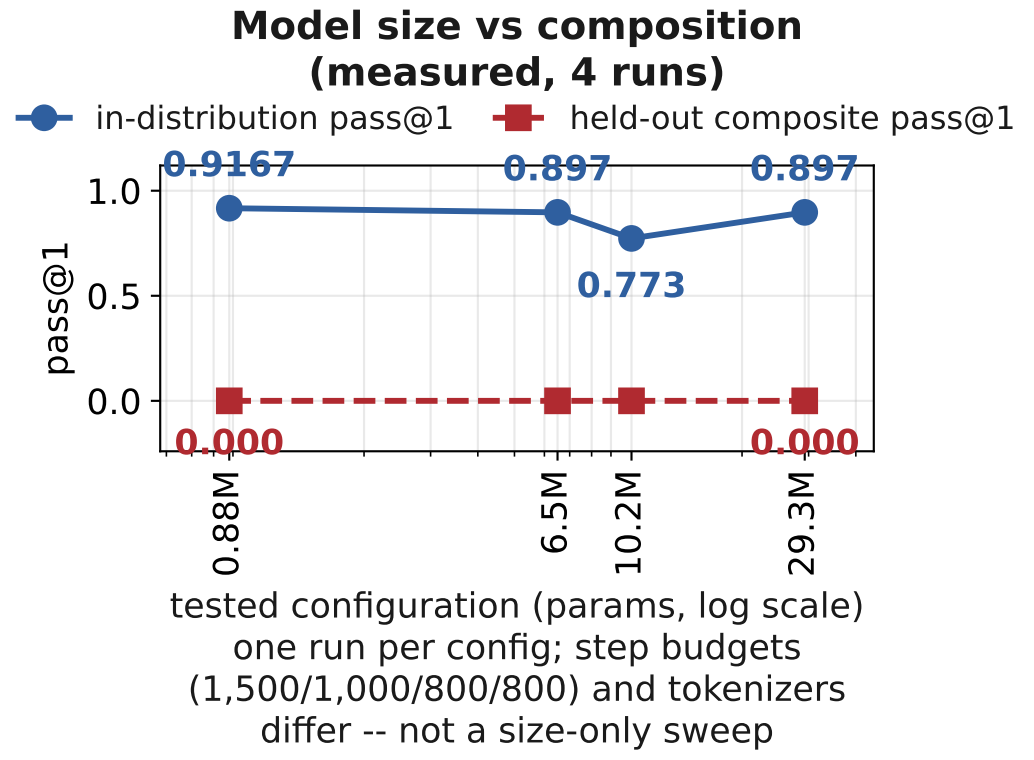
Model size vs composition
(measured, 4 runs)
in-distribution pass@1
held-out composite pass@1
pass@1
1.0
0.5
0.0
0.9167
0.897
0.897
0.773
0.000
0.000
0.88M
6.5M
10.2M
29.3M
tested configuration (params, log scale)
one run per config; step budgets
(1,500/1,000/800/800) and tokenizers
differ -- not a size-only sweep