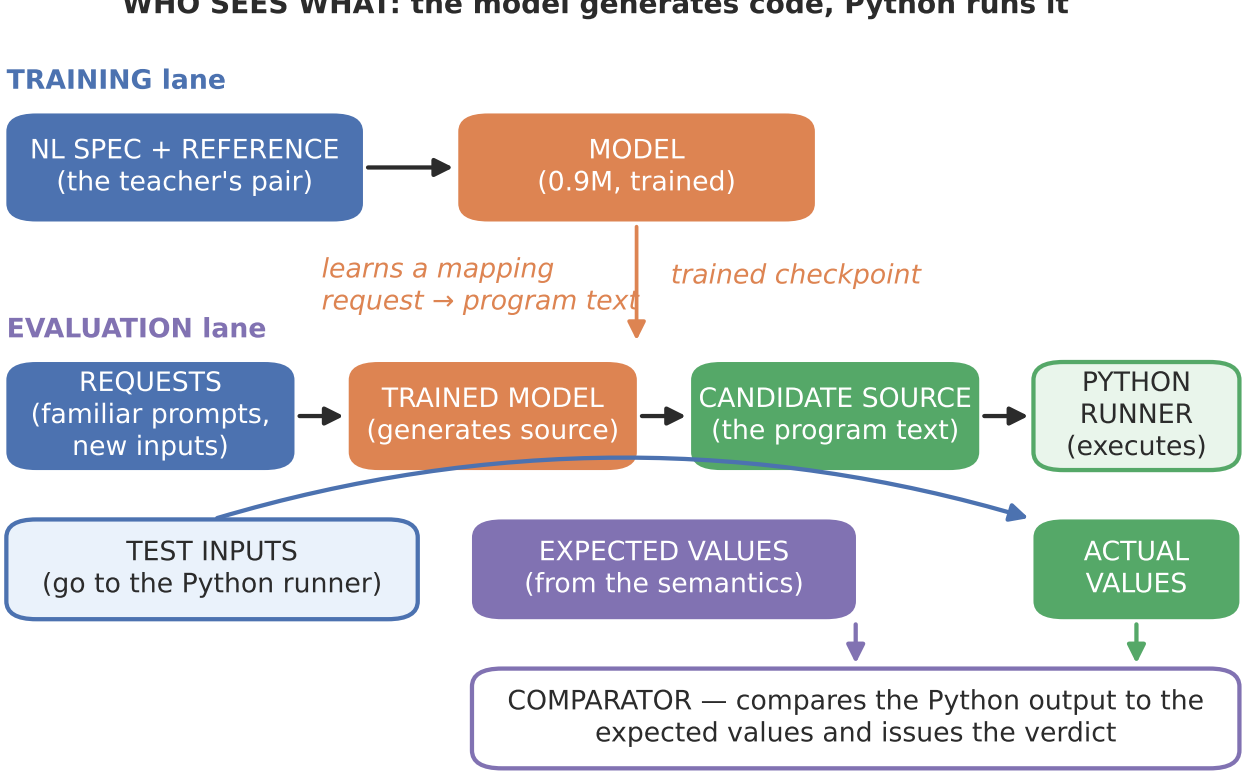
WHO SEES WHAT: the model generates code, Python runs it
TRAINING lane
NL SPEC + REFERENCE
(the teacher's pair)
MODEL
(0.9M, trained)
learns a mapping
request → program text
trained checkpoint
EVALUATION lane
REQUESTS
(familiar prompts,
new inputs)
TRAINED MODEL
(generates source)
CANDIDATE SOURCE
(the program text)
PYTHON
RUNNER
(executes)
TEST INPUTS
(go to the Python runner)
EXPECTED VALUES
(from the semantics)
ACTUAL
VALUES
COMPARATOR — compares the Python output to the
expected values and issues the verdict
CANDIDATE EXECUTION · PURE FUNCTION (no I/O, no network, no random) · AST-WHITELISTED