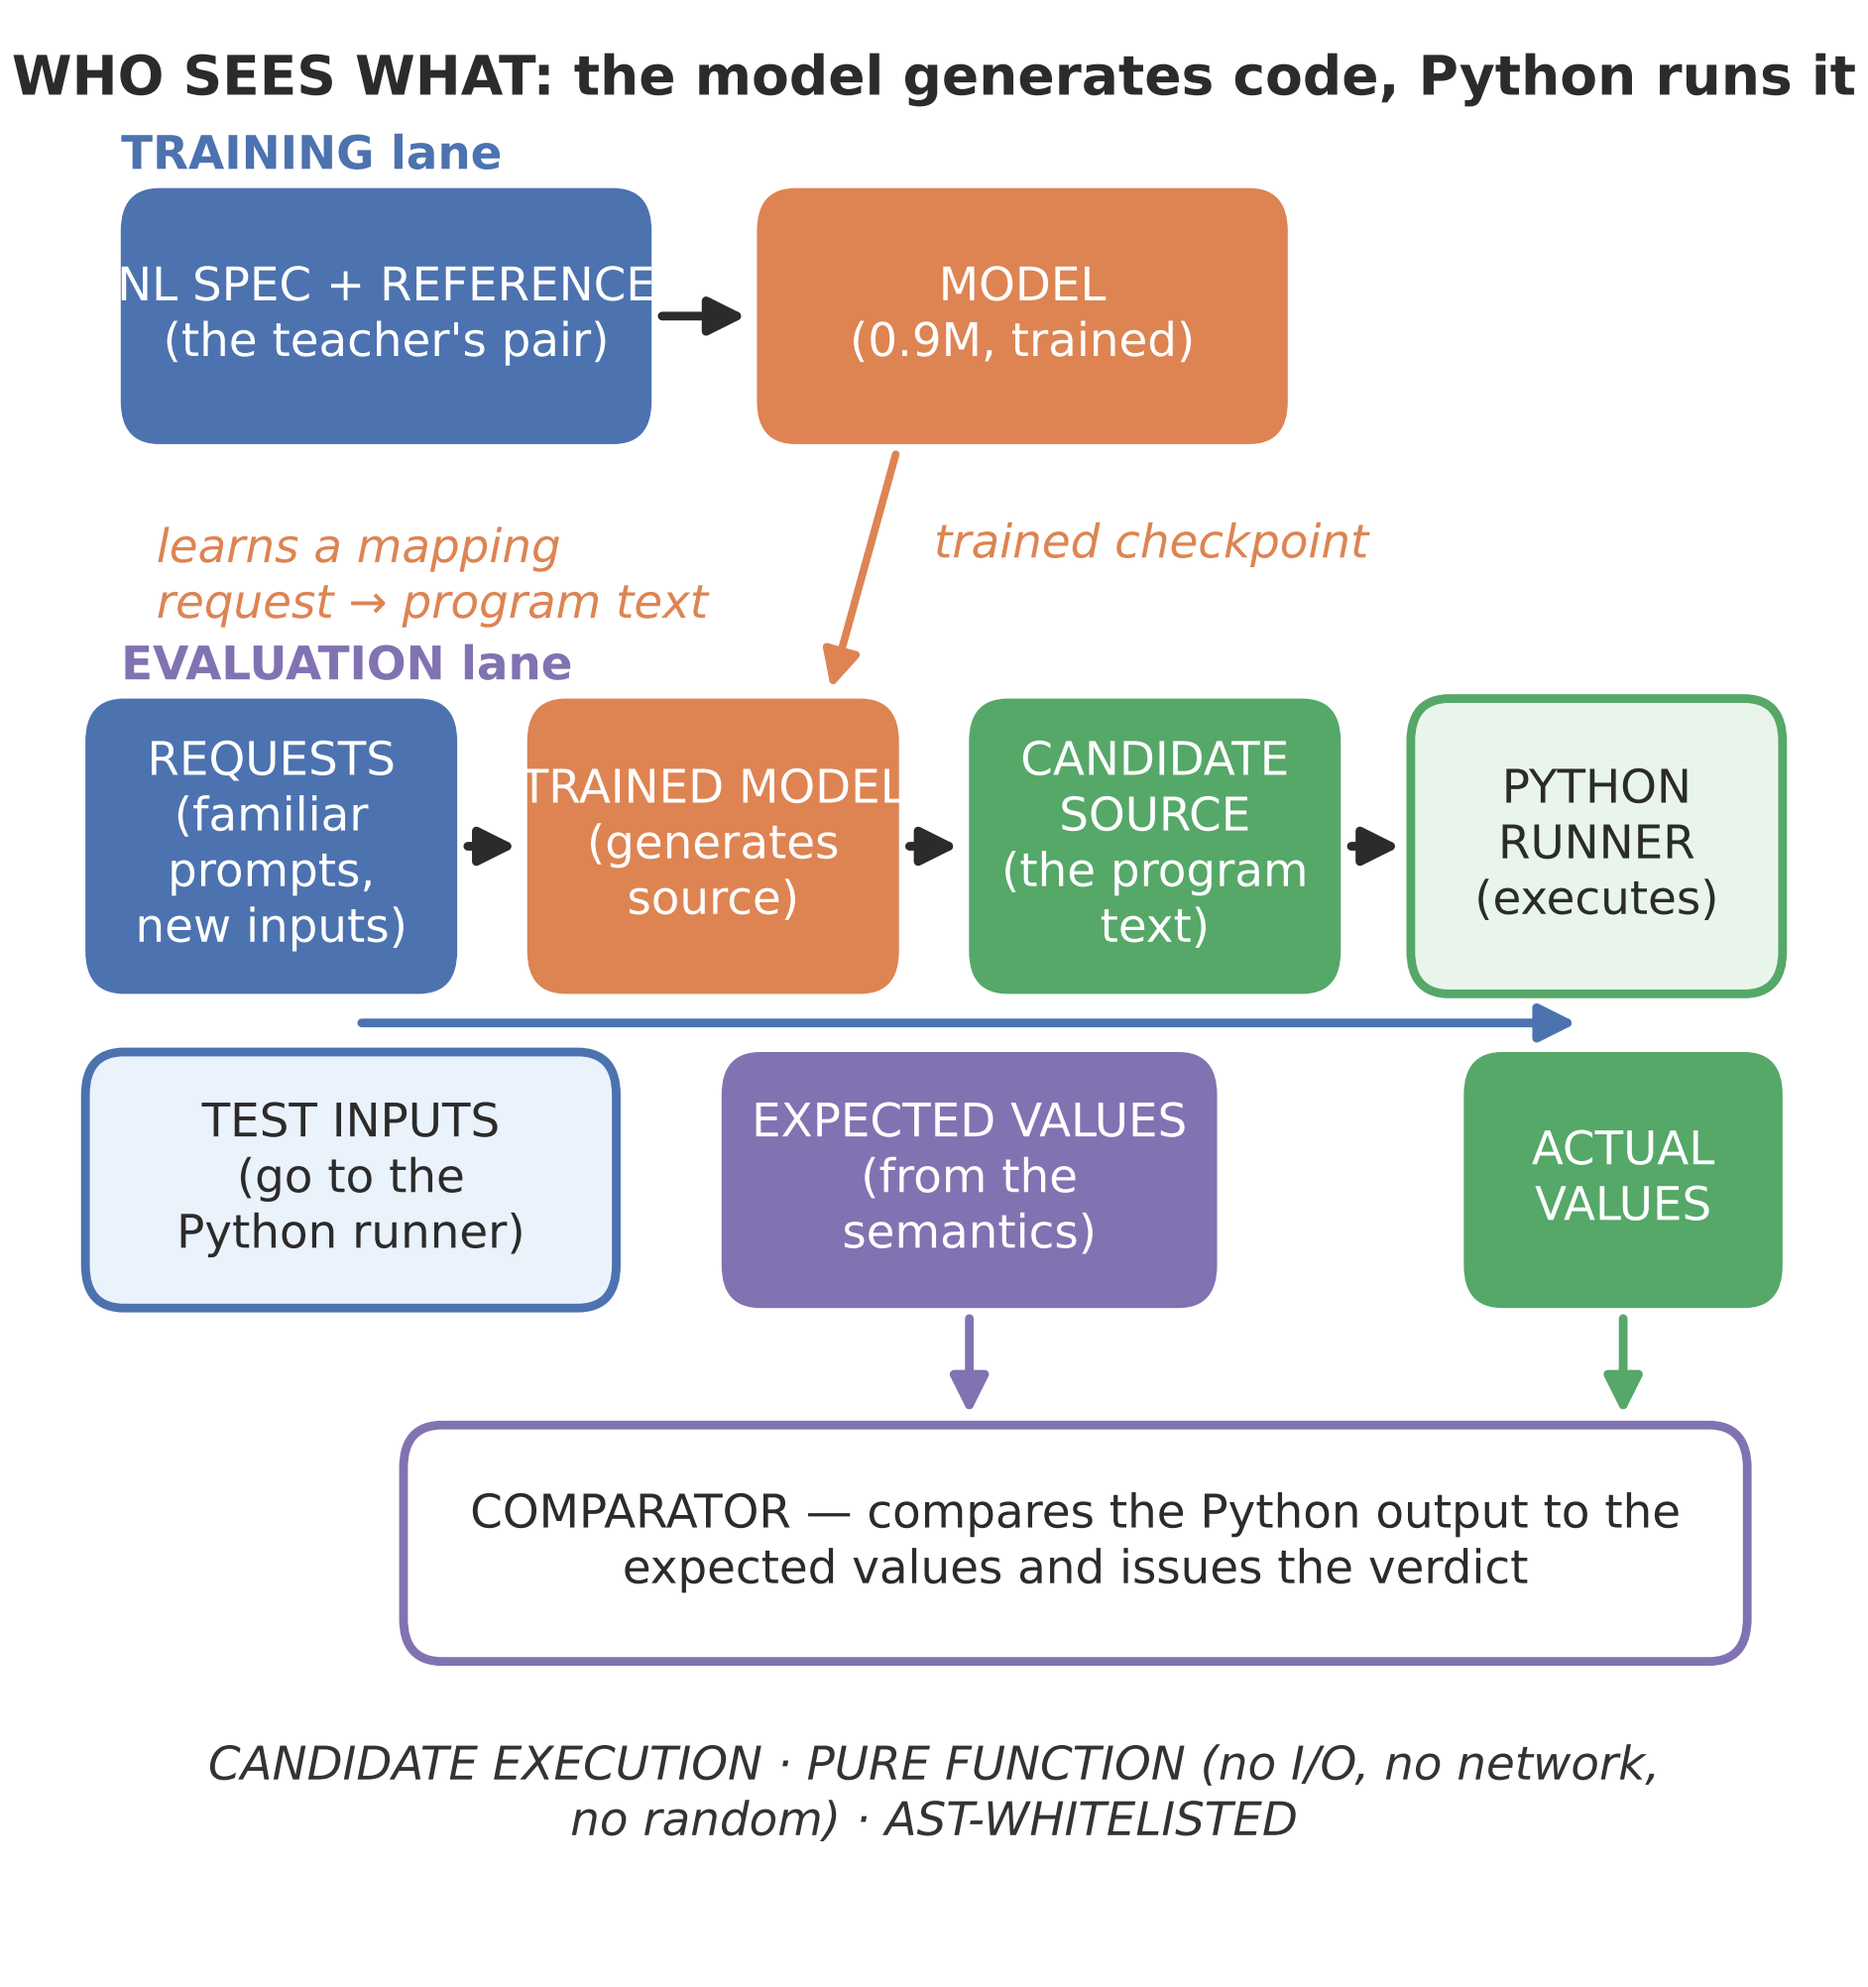
WHO SEES WHAT: the model generates code, Python runs it
TRAINING lane
NL SPEC + REFERENCE
(the teacher's pair)
MODEL
(0.9M, trained)
learns a mapping
request → program text
trained checkpoint
EVALUATION lane
REQUESTS
(familiar
prompts,
new inputs)
TRAINED MODEL
(generates
source)
CANDIDATE
SOURCE
(the program
text)
PYTHON
RUNNER
(executes)
TEST INPUTS
(go to the
Python runner)
EXPECTED VALUES
(from the
semantics)
ACTUAL
VALUES
COMPARATOR — compares the Python output to the
expected values and issues the verdict
CANDIDATE EXECUTION · PURE FUNCTION (no I/O, no network,
no random) · AST-WHITELISTED